



WEAVEHACKS 4 · APPLIED INTELLIGENCE

CCIP

Content Creators Intelligence Platform — multi-agent HumaneBench grading for YouTube & X/Twitter communities.

LangGraph

WandB Weave

Redis VL

CopilotKit

Apify

HumaneBench v3.0

THE PROBLEM

Creators are flying blind on community health.

YouTube & X creators get thousands of comments. Most are noise — spam, toxicity, low-effort reactions. The few high-signal comments that represent real community sentiment are buried. There is no objective, repeatable way to grade them.

95%

COMMENTS NEVER READ BY CREATOR

0

STANDARD FRAMEWORKS FOR COMMUNITY
HEALTH


∞

SUBJECTIVITY IN MODERATION DECISIONS

THE SOLUTION

One URL → full community intelligence in seconds

Paste any YouTube or X link. CCIP's LangGraph supervisor orchestrates 4 specialist agents — scraping content, generating a vibe blueprint, indexing comments into Redis VL, and grading every comment with HumaneBench v3.0. All traced in WandB Weave.



THE FRAMEWORK

HumaneBench v3.0

8 principles. Objective scores. A–F grades. Every comment evaluated across Respect
Attention, Meaningful Choices, Enhance Capabilities, Dignity & Safety, Healthy Relationships,
Long-term Wellbeing, Transparency, and Equity.



Attention

Deserves viewer time



Choices

Offers real options



Growth

Helps audience grow



Safety

Respectful & safe



Relationships

Positive community



Wellbeing

Lasting interests



Honesty

Authentic & truthful



Equity

Inclusive & fair

GRADING SYSTEM

+1.0 to −1.0 per principle. A+ to F.

Each of 8 principles scores +1.0 (Exemplary) / +0.5 (Acceptable) / −0.5 (Concerning) / −1.0 (Violation). Final score is the mean — grades map A+ (≥ 0.875) down to F.

GRADE	SCORE RANGE	WHAT IT MEANS
A+	≥ 0.875	Exemplary across all principles
A	≥ 0.625	Strong, high-signal community contribution
B	≥ 0.375	Good — minor room for improvement
C	≥ 0.125	Average — some concerning elements
D	≥ -0.125	Below average — notable violations
F	< -0.125	Harmful — violates multiple principles

Built on W&B + LangGraph + Redis



LangGraph Supervisor

Orchestrates the 4-agent pipeline. State machine with typed edges — scrape → blueprint → index → grade → done. Each agent is a LangGraph node with weave tracing.



WandB Weave + LLM

All agent calls traced via @weave.op. LLM inference through WandB's OpenAI-compatible endpoint — Qwen3-14B-Instruct for blueprint + grading.



Redis VL + RAG

Comments embedded with sentence-transformers (384-dim HNSW). HumaneBench KB stored in Redis VL. Grader uses RAG to pull relevant principle examples.



CopilotKit AI Chat

Three in-context AI panels — Blueprint AI, Vibe Audit AI, Prompt Chip AI — each wired to relevant CopilotReadable context. Powered by WandB LLM.

WHAT YOU GET

Three intelligence layers, one pipeline.



Vibe Blueprint

Creator's emotional state, community need (true intent), interaction boundaries, and AI-generated contextual prompt chips for fans.



Comment Audit Feed

Every comment graded A+ → F with per-principle dot breakdown. Click any comment to open the AI panel for deep-dive analysis.



Vibe Report

Monetizable Retention State — community alignment %, tier depth, high-signal ratio, grade distribution, and recommended action.

Every agent call is a Weave trace.

CCIP uses WandB Weave as the observability backbone. Every `@weave.op`-decorated specialist — scrape, blueprint generation, embedding, grading — appears as a trace in your W&B project. You get full cost, latency, and output lineage for every run.



Full trace lineage

Parent run → child agents → LLM calls. Click into any comment grade to see the exact prompt + response chain.



WandB LLM inference

OpenAI-compatible endpoint. One API key, Qwen3-14B for both blueprint and grading — all usage tracked in W&B.



Zero extra infra

WandB handles model serving and tracing. We bring LangGraph + Redis VL. No separate model infrastructure needed.

LIVE DEMO



Let's run it.

Drop a YouTube URL. Watch 4 agents fire in sequence. See every comment graded A+ → F in real time. Ask the AI anything about the results.

Scrape → Blueprint → Index → Grade

~60s end-to-end

Weave traces live



WEAVEHACKS 4

Content Creators Intelligence Platform

LangGraph · WandB Weave · Redis VL · CopilotKit · HumaneBench
v3.0

Built by **Abhijit Betigeri** · abhijitbetigeri29@gmail.com

github.com/abhijitbetigeri/WB-Hack